

Disk Encryption and Message Authentication

Applied Cryptography – Spring 2024


Bart Mennink

January 31, 2024

Institute for Computing and Information Sciences
Radboud University

1 / 25

Recap: Block Ciphers



- Using key K , message M is bijectively transformed to ciphertext C
- Key, plaintext, and ciphertext are typically of fixed size
- Example [DR01]:

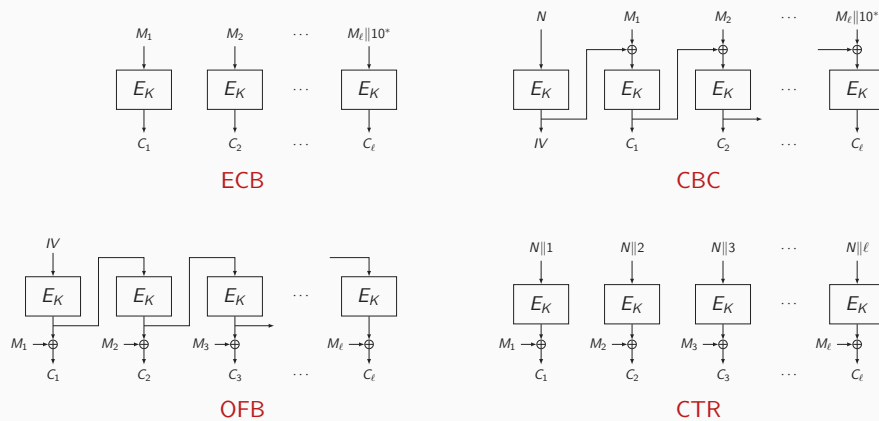
$$\text{AES-128: } \{0, 1\}^{128} \times \{0, 1\}^{128} \rightarrow \{0, 1\}^{128}$$

$$(K, M) \mapsto C$$

- For fixed key, E_K is invertible and the inverse is denoted as E_K^{-1}
- A good block cipher should behave like a random permutation

2 / 25

Block Cipher Encryption Modes



3 / 25

Last Lecture

- Keyed cryptographic constructions are **modular**:
 - A small primitive is turned into a larger mode of use
 - Typically, however, we only know how to build primitives that behave like **random permutations**
 - Most notable, **block ciphers** like AES
 - Mostly historically, but people still use them **a lot!**
 - E.g., each website over HTTPS sets up a TLS connection and reportedly over 70% over these connections use AES-GCM
- In this lecture:
 - Disk encryption
 - Message authentication
 - Beginning of authenticated encryption

4 / 25

Outline

Disk Encryption

Message Authentication

Intermezzo: Universal Hashing

Example: Wegman-Carter(-Shoup) and Protected Hash


Example: CBC-MAC

Authenticated Encryption (Teaser)

5 / 25

Disk Encryption

Disk Encryption



Main Goals


- Confidentiality
- Efficient in encryption and decryption
- No ciphertext expansion
- User friendly
- Incrementality ... but not too much

These Slides

- High-level idea
- Core behind VeraCrypt, open source disk encryption tool

6 / 25

VeraCrypt Disk Encryption (1/2)



Encryption Scheme

- Actually encrypts all your data
- Requires some symmetric key ... but we only have our password???

KDF – Key Derivation Function

- Can “generate” a key from “limited information”
- But is this secure?

PRNG – Pseudorandom Number Generator


- Accumulates entropy from, e.g., mouse movements
- Turns it into a random looking string, the **salt**
- Salt fed to KDF to prevent **dictionary attacks**

7 / 25

VeraCrypt Disk Encryption (2/2)

In VeraCrypt


- Encryption scheme
 - XTS with AES-128 (more in Lecture 3)
 - Widely used standard for data encryption
- KDF
 - PBKDF2, hoped to behave like a pseudorandom function (more in Lecture 4)
 - Salt is 512 bits, so input has quite some "randomness"
- Weakest spot?
 - The user! Short passwords are easy to guess
 - Salvaged by adjusting the number of rounds in PBKDF2



8 / 25

Message Authentication


Message Authentication (1/2)



- Using key K , message M is signed with tag T
- Associated verification function takes K and (M, T) and outputs
 - \top if tag is correct
 - \perp if tag is incorrect
- Key and tag are typically of fixed size
- Message is typically of arbitrary length
- Sometimes, additional nonce: MAC-evaluations should be for unique nonce

9 / 25

Message Authentication (2/2)



- Applications:
 - Message authentication: append tag to message
 - Entity authentication: compute tag over challenge
- Security goal: MAC_K should behave like a random function
 - I.e., $\text{Adv}_{\text{MAC}}^{\text{prf}}(q)$ should be small
- Often, one adopts a weaker notion, called unforgeability


10 / 25

Message Authentication Security

- MAC security defined using **unforgeability**
- Adversary \mathcal{A} has access to MAC_K and VFY_K
- It can query both oracles interchangeably
 - No** query VFY_K for an output of MAC_K
 - Optional: **no** repeated N for MAC_K
- \mathcal{A} mounts a **forgery** if VFY_K ever outputs \top
- Its advantage is defined as:

$$\text{Adv}_{\text{MAC}}^{\text{unf}}(\mathcal{A}) = \Pr(\mathcal{A}^{\text{MAC}_K, \text{VFY}_K} \text{ forges})$$

- $\text{Adv}_{\text{MAC}}^{\text{unf}}(q_m, q_v)$: supremal advantage over any \mathcal{A} with:
 - query complexity q_m to MAC_K
 - query complexity q_v to VFY_K



11 / 25

Pseudorandomness Implies Unforgeability

- Consider a MAC function MAC_K
- Intuitively, if outputs of MAC_K look random, they should also be hard to forge
- Forging would then correspond to "guessing" a random output of MAC_K
- There is a well-established result that proves this:
 - Original result dates back to 1984, Goldreich et al. [GGM84]
 - In currently established formalism: Bellare et al. [BKR94, BGM04]

$$\text{Adv}_{\text{MAC}}^{\text{unf}}(q_m, q_v) \leq \text{Adv}_{\text{MAC}}^{\text{prf}}(q_m + q_v) + \frac{q_v}{2^t}$$

- Proof: see Theorem 6.2.2 of "Intro2Crypto-symmetric.pdf"

12 / 25


Universal Hashing

Universal Hash Functions

- Same interface as MAC, but weaker security requirements
 - Con**: function should not be **exposed** to attacker
 - Pro**: function can often be evaluated **much cleanly**
- Hash function family $H : \{0, 1\}^k \times \{0, 1\}^* \rightarrow \{0, 1\}^t$ is called
 - δ -**universal** if $\Pr_K(H_K(M) = H_K(M')) \leq \delta \quad (\forall M \neq M')$
 - ϵ -**XOR-universal** if $\Pr_K(H_K(M) \oplus H_K(M') = T) \leq \epsilon \quad (\forall M \neq M', T)$


GHASH

- Addition and multiplication over finite field
- $\ell 2^{-t}$ -(XOR-)universal [MV04]



13 / 25


Wegman-Carter



- Introduced in 1986 [WC81]
- Process arbitrary length M through universal hash, mask with $F_K(N)$
- Secure MAC function if
 - H is ϵ -XOR-universal
 - F is pseudorandom function
 - Nonce N **never reused** in MAC-queries
- $\text{Adv}_{\text{WC}}^{\text{unf}}(q_m, q_v) \leq q_v \epsilon + \text{Adv}_F^{\text{prf}}(q_m + q_v)$ [Ber05]

14 / 25

Wegman-Carter-Shoup



- PRFs hard to construct, PRPs easy: **take a block cipher E**
- Described by Shoup in 1996 [Sho96]
- Secure MAC function if
 - H is ε -XOR-universal
 - E is pseudorandom permutation
 - Nonce N never reused
- $\text{Adv}_{\text{WC}}^{\text{unf}}(q_m, q_v) \leq e^{(q_m+1)q_v/2^n} q_v \varepsilon + \text{Adv}_E^{\text{PRP}}(q_m + q_v)$ [Ber05]

15 / 25

Poly1305 [Ber05]

- Wegman-Carter-Shoup construction with **cleverly chosen** universal hash
- Originally defined for AES, but commonly used with ChaCha20
- Observation: 128-bit string X can be written as integer in $[0, 2^{128} - 1]$:

$$X[0] \| X[1] \| \dots \| X[127] \in \{0, 1\}^{128} \sim X[0] + 2X[1] + \dots + 2^{127}X[127] \in [0, 2^{128} - 1]$$
- Keys $K, R \in \{0, 1\}^{128}$ with 22 bits of R fixed
- On input of nonce $N \in \{0, 1\}^{128}$ and message $M \in \{0, 1\}^*$:
 - M is partitioned into ℓ 128-bit blocks M_i
 - Each block (except last one) is appended with a 1
 - Call the updated blocks C_1, \dots, C_ℓ
 - Note, $C_1, \dots, C_{\ell-1}$ are integers in $[2^{128}, 2^{129} - 1]$
 - $T = \left(((C_1 \cdot R^\ell + C_2 \cdot R^{\ell-1} + \dots + C_\ell \cdot R) \bmod 2^{130} - 5) + E_K(N) \right) \bmod 2^{128}$

16 / 25


Protected Hash (1/2)



- Not a dedicated construction “as such”, but appears quite frequently in disguise
 - CBC-MAC [BKR94]
 - Protected counter sum [Ber99]
- Process arbitrary length M through universal hash, protect with $F_K(\cdot)$ or $E_K(\cdot)$

17 / 25


Protected Hash (2/2)



- Secure MAC function if
 - H is δ -universal
 - F is pseudorandom function
- $\text{Adv}_{\text{WC}}^{\text{unf}}(q_m, q_v) \leq \binom{q_m+q_v}{2} \delta + q_v/2^t + \text{Adv}_F^{\text{PRF}}(q_m + q_v)$ [Sho04]
- In case we protect with $E_K(N)$, extra loss of $\binom{q_m+q_v}{2}/2^n$

18 / 25


Cipher Block Chaining MAC (CBC-MAC) Mode



- In CBC encryption: C_i depends on M_1, \dots, M_i
- Idea for message authentication:
 - Apply CBC with $IV = 0$ to padded message M
 - Define tag T to be the last ciphertext block
 - Important: **discard all other ciphertext blocks!**
- Turns out to be secure if messages are **prefix-free**

19 / 25

Cipher Block Chaining MAC (CBC-MAC) Mode: Weakness




- In general, CBC-MAC can be distinguished from random in two queries:
 - Query M_1 , tag equals $T = E_K(M_1)$
 - Query $M_1 \| (M_1 \oplus T)$, tag equals

$$E_K(E_K(M_1) \oplus M_1 \oplus T) = E_K(T \oplus M_1 \oplus T) = E_K(M_1) = T$$
- A random oracle would give independent responses
- CBC-MAC is **not** unforgeable and **definitely not** PRF-secure
- Note: attack ignores padding, but this can be dealt with

20 / 25

Cipher Block Chaining MAC (CBC-MAC) Mode: Fix




- Length-extension attack can be resolved by "special" finalization
- Solution 1: **mask** last block with dedicated key L (known as C-MAC)
- Solution 2: apply **independent** last primitive call
 - Can be seen as protected hash construction
- Both constructions indistinguishable from RO up to around $(\frac{q}{2})/2^n$

21 / 25

Authenticated Encryption (Teaser)


Authenticated Encryption



- Using key K :
 - Message M is encrypted in ciphertext C
 - Associated data A and message M are authenticated using T
- Nonce N randomizes the scheme
- Key, nonce, and tag are typically of fixed size
- Associated data, message, and ciphertext could be arbitrary length

22 / 25


Authenticated Decryption



- Authenticated decryption needs to satisfy that
 - Message disclosed if tag is **correct**
 - Message is not leaked if tag is **incorrect**
- Correctness: $AE_K^{-1}(N, A, AE_K(N, A, M)) = M$

23 / 25

Authenticated Encryption Security



- Two oracles: (AE_K, AE_K^{-1}) (for secret key K) and $(\$, \perp)$ (secret)
- Distinguisher \mathcal{D} has query access to one of these
 - unique nonce for each encryption query, and no trivial queries
- \mathcal{D} tries to determine which oracle it communicates with

- Its advantage is defined as:


$$\text{Adv}_{\text{AE}}^{\text{ae}}(\mathcal{D}) = \Delta_{\mathcal{D}}(AE_K, AE_K^{-1}; \$, \perp) = \left| \Pr(\mathcal{D}^{AE_K, AE_K^{-1}} = 1) - \Pr(\mathcal{D}^{\$, \perp} = 1) \right|$$

- $\text{Adv}_{\text{AE}}^{\text{ae}}(q_e, q_v)$: supramal advantage over any \mathcal{D} with query complexity q_e, q_v

24 / 25

Generic Composition

- Generic constructions for AE:
 - $\text{Enc} + \text{MAC} = \text{AE}$
- Bellare and Namprempe (2000): 3 basic approaches (N to both Enc/MAC , A to MAC)



- Used in SSH
- Generically insecure
 - $\text{MAC}_L(M) = M \parallel T?$
- Used in TLS
- Mildly insecure
- Padding oracle attack
- Used in IPSec
- Most secure variant
- Ciphertext integrity

25 / 25