



Authenticated Encryption

Applied Cryptography – Spring 2024

Bart Mennink

February 19, 2024

Institute for Computing and Information Sciences
Radboud University

Encryption


- Security goal: confidentiality
- Examples: ECB, counter mode

Authentication

- Security goal: data integrity
- Examples: CBC-MAC, Poly1305


Authenticated encryption combines both

Authenticated Encryption




- Using key K :
 - Message M is encrypted in ciphertext C
 - Associated data A and message M are authenticated using T
- Nonce N randomizes the scheme

Authenticated Encryption




- Using key K :
 - Message M is encrypted in ciphertext C
 - Associated data A and message M are authenticated using T
- Nonce N randomizes the scheme
- Authenticated decryption discloses M if and only if T is **correct**

Authenticated Encryption Security



- Two oracles: (AE_K, AE_K^{-1}) (for secret key K) and $(\$, \perp)$ (secret)
- Distinguisher \mathcal{D} has query access to one of these
→ unique nonce for each encryption query, and no trivial queries
- \mathcal{D} tries to determine which oracle it communicates with

Authenticated Encryption Security



- Two oracles: (AE_K, AE_K^{-1}) (for secret key K) and $(\$, \perp)$ (secret)
- Distinguisher \mathcal{D} has query access to one of these
→ unique nonce for each encryption query, and no trivial queries
- \mathcal{D} tries to determine which oracle it communicates with

- Its advantage is defined as:

$$\mathbf{Adv}_{AE}^{\text{ae}}(\mathcal{D}) = \Delta_{\mathcal{D}}(AE_K, AE_K^{-1}; \$, \perp) = \left| \Pr(\mathcal{D}^{AE_K, AE_K^{-1}} = 1) - \Pr(\mathcal{D}^{\$, \perp} = 1) \right|$$

- $\mathbf{Adv}_{AE}^{\text{ae}}(q_e, q_v)$: supremal advantage over any \mathcal{D} with query complexity q_e, q_v

Authenticated Encryption Design

- Simple Example

- Example: GCM Authenticated Encryption

- Role of the Nonce, and GCM-SIV Authenticated Encryption

Tweakable Block Ciphers

- Example: Θ CB Authenticated Encryption


Building Tweakable Block Ciphers

- Application to Authenticated Encryption

Authenticated Encryption Design


Encryption

- Input: (N, M)




Encryption

- Input: (N, M)
- Compute keystream $Z_1 \parallel Z_2$




Simple Example



Encryption

- Input: (N, M)
- Compute keystream $Z_1 \parallel Z_2$
- Output:
 - $C = Z_1 \oplus M$


Simple Example



Encryption

- Input: (N, M)
- Compute keystream $Z_1 \parallel Z_2$
- Output:
 - $C = Z_1 \oplus M$
 - $T = Z_2 \oplus (M \otimes L)$

Simple Example




Encryption

- Input: (N, M)
- Compute keystream $Z_1 \parallel Z_2$
- Output:
 - $C = Z_1 \oplus M$
 - $T = Z_2 \oplus (M \otimes L)$

Decryption

- Input: (N, C, T)

Simple Example




Encryption

- Input: (N, M)
- Compute keystream $Z_1 \parallel Z_2$
- Output:
 - $C = Z_1 \oplus M$
 - $T = Z_2 \oplus (M \otimes L)$

Decryption

- Input: (N, C, T)
- Compute keystream $Z_1 \parallel Z_2$
- Compute $M = Z_1 \oplus C$

Simple Example




Encryption

- Input: (N, M)
- Compute keystream $Z_1 \parallel Z_2$
- Output:
 - $C = Z_1 \oplus M$
 - $T = Z_2 \oplus (M \otimes L)$

Decryption

- Input: (N, C, T)
- Compute keystream $Z_1 \parallel Z_2$
- Compute $M = Z_1 \oplus C$
- Compute $T^* = Z_2 \oplus (M \otimes L)$

Simple Example




Encryption

- Input: (N, M)
- Compute keystream $Z_1 \parallel Z_2$
- Output:
 - $C = Z_1 \oplus M$
 - $T = Z_2 \oplus (M \otimes L)$

Decryption

- Input: (N, C, T)
- Compute keystream $Z_1 \parallel Z_2$
- Compute $M = Z_1 \oplus C$
- Compute $T^* = Z_2 \oplus (M \otimes L)$
- Output: $\begin{cases} M & \text{if } T = T^* \\ \perp & \text{otherwise} \end{cases}$


Simple Example: Confidentiality



Confidentiality

- Consider new query (N, M)
- N should be **fresh**


Simple Example: Confidentiality



Confidentiality

- Consider new query (N, M)
- N should be **fresh**
- Random $Z_1 \parallel Z_2$
(if F is a good stream cipher)

Simple Example: Confidentiality




Confidentiality

- Consider new query (N, M)
- N should be **fresh**
- Random $Z_1 \parallel Z_2$
(if F is a good stream cipher)
- Random (C, T)


Simple Example: Authenticity

Authenticity

- Consider forgery attempt (N, C, T)




Simple Example: Authenticity



Authenticity

- Consider forgery attempt (N, C, T)
- N could be **repeated** nonce


Simple Example: Authenticity



Authenticity

- Consider forgery attempt (N, C, T)
- N could be **repeated** nonce
- N fresh:
 - T^* is random, unpredictable


Simple Example: Authenticity



Authenticity

- Consider forgery attempt (N, C, T)
- N could be **repeated** nonce
- N fresh:
 - T^* is random, unpredictable
- N repeated:
 - Let (N, M', C', T') be old


Simple Example: Authenticity



Authenticity

- Consider forgery attempt (N, C, T)
- N could be **repeated** nonce
- N fresh:
 - T^* is random, unpredictable
- N repeated:
 - Let (N, M', C', T') be old
 - $M = Z_1 \oplus C = M' \oplus C' \oplus C$


Simple Example: Authenticity



Authenticity

- Consider forgery attempt (N, C, T)
- N could be **repeated** nonce
- N fresh:
 - T^* is random, unpredictable
- N repeated:
 - Let (N, M', C', T') be old
 - $M = Z_1 \oplus C = M' \oplus C' \oplus C$
 - $T^* = Z_2 \oplus (M \otimes L)$
 - $T^* = T' \oplus ((M \oplus M') \otimes L)$
 - $T^* = T' \oplus ((C \oplus C') \otimes L)$


Simple Example: Authenticity



Authenticity

- Consider forgery attempt (N, C, T)
- N could be **repeated** nonce
- N fresh:
 - T^* is random, unpredictable
- N repeated:
 - Let (N, M', C', T') be old
 - $M = Z_1 \oplus C = M' \oplus C' \oplus C$
 - $T^* = Z_2 \oplus (M \otimes L)$
 - $T^* = T' \oplus ((M \oplus M') \otimes L)$
 - $T^* = T' \oplus ((C \oplus C') \otimes L)$
 - Forgery successful if
$$T \oplus T' = (C \oplus C') \otimes L$$

Simple Example: Authenticity




Authenticity


- Consider forgery attempt (N, C, T)
- N could be **repeated** nonce
- N fresh:
 - T^* is random, unpredictable
- N repeated:
 - Let (N, M', C', T') be old
 - $M = Z_1 \oplus C = M' \oplus C' \oplus C$
 - $T^* = Z_2 \oplus (M \otimes L)$
 - $T^* = T' \oplus ((M \oplus M') \otimes L)$
 - $T^* = T' \oplus ((C \oplus C') \otimes L)$
 - Forgery successful if
 - $T \oplus T' = (C \oplus C') \otimes L$
 - Requires **guessing** L

Simple Example: How to Support Variable Length?

Suppose M is Variable-Length?




Simple Example: How to Support Variable Length?



Suppose M is Variable-Length?

- Output F should be twice as large?
- \otimes over arbitrary # of bits?


Simple Example: How to Support Variable Length?



Suppose M is Variable-Length?

- Output F should be twice as large?
- \otimes over arbitrary # of bits?
- $|M| + n$ bits turns out to suffice:
 - Use streaming mode for F
 - Replace $M \otimes L$ by $H_L(M)$

Simple Example: How to Support Variable Length?




Suppose M is Variable-Length?

- Output F should be twice as large?
- \otimes over arbitrary # of bits?
- $|M| + n$ bits turns out to suffice:
 - Use streaming mode for F
 - Replace $M \otimes L$ by $H_L(M)$

What about AD A?

Simple Example: How to Support Variable Length?




Suppose M is Variable-Length?

- Output F should be twice as large?
- \otimes over arbitrary # of bits?
- $|M| + n$ bits turns out to suffice:
 - Use streaming mode for F
 - Replace $M \otimes L$ by $H_L(M)$

What about AD A?

- Can be processed by H_L as well:
 - $H_L(A, M)$

Simple Example: How to Support Variable Length?



Suppose M is Variable-Length?


- Output F should be twice as large?
- \otimes over arbitrary # of bits?
- $|M| + n$ bits turns out to suffice:
 - Use streaming mode for F
 - Replace $M \otimes L$ by $H_L(M)$

What about AD A ?

- Can be processed by H_L as well:
 - $H_L(A, M)$

This is almost exactly GCM!

Simple Example: How to Support Variable Length?



Suppose M is Variable-Length?

- Output F should be twice as large?
- \otimes over arbitrary $\#$ of bits?
- $|M| + n$ bits turns out to suffice:
 - Use streaming mode for F
 - Replace $M \otimes L$ by $H_L(M)$

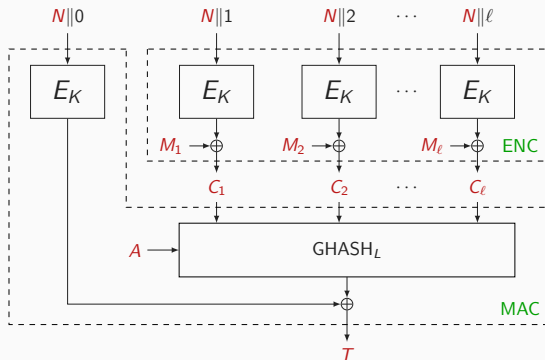
What about AD A?

- Can be processed by H_L as well:
 - $H_L(A, M)$

This is almost exactly GCM!

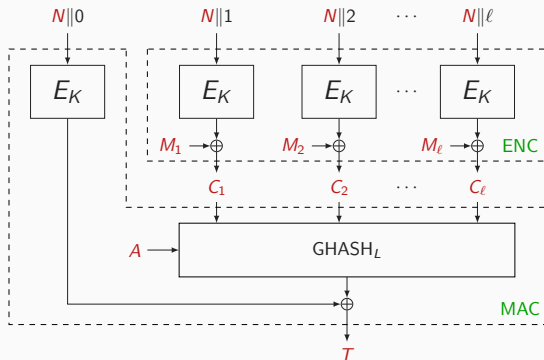
- Encrypt-then-MAC: $H_L(A, C)$
- Take CTR mode for F

GCM for 96-bit Nonce N




- McGrew and Viega (2004)
- EtM design
- Widely used (TLS!)
- Patent-free

GCM for 96-bit Nonce N



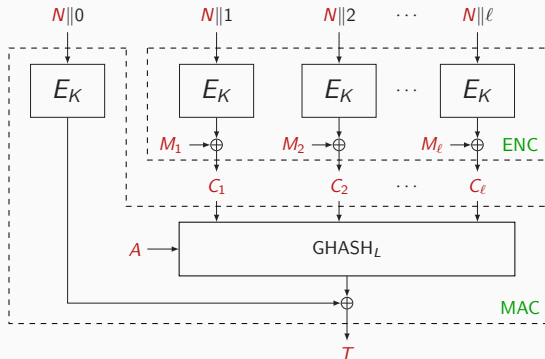
- McGrew and Viega (2004)
- EtM design
- Widely used (TLS!)
- Patent-free
- Parallelizable
- Evaluates E only (no E^{-1})
- Provably secure (if E is PRP)
- Very efficient in HW
- Reasonably efficient in SW

GCM for 96-bit Nonce N



- McGrew and Viega (2004)
- EtM design
- Widely used (TLS!)
- Patent-free
- Parallelizable
- Evaluates E only (no E^{-1})
- Provably secure (if E is PRP)
- Very efficient in HW
- Reasonably efficient in SW
- Note: equally popular is ChaCha20-Poly1305!



GCM for 96-bit Nonce N



- McGrew and Viega (2004)
- EtM design
- Widely used (TLS!)
- Patent-free
- Parallelizable
- Evaluates E only (no E^{-1})
- Provably secure (if E is PRP)
- Very efficient in HW
- Reasonably efficient in SW
- Note: equally popular is ChaCha20-Poly1305!



What happens if nonce is re-used?

Nonce = "Number Used Once"



- Nonces N and N' should be distinct for two different evaluations


Nonce = "Number Used Once"




- Nonces N and N' should be distinct for two different evaluations
- What happens if a nonce would be repeated?

What if a Nonce is Repeated?

M




M'




- Consider evaluations for identical nonce but distinct M, M'

What if a Nonce is Repeated?

M





M'





- Consider evaluations for identical nonce but distinct M, M'
- Key streams will be identical

What if a Nonce is Repeated?




- Consider evaluations for identical nonce but distinct M, M'
- Key streams will be identical
- Ciphertexts satisfy $C \oplus C' = M \oplus M' \rightarrow$ attacker knew C' in advance

What if a Nonce is Repeated?




- Consider evaluations for identical nonce but distinct M, M'
- Key streams will be identical
- Ciphertexts satisfy $C \oplus C' = M \oplus M' \rightarrow$ attacker knew C' in advance
- Tags satisfy $T \oplus T' = M \otimes L \oplus M' \otimes L = (M \oplus M') \otimes L \rightarrow$ key recovery

Guaranteeing Uniqueness of Nonce




Guaranteeing Uniqueness of Nonce



- Issues with nonce generation:
 - Counter needs storage
 - Need synchronization or transmission
 - Efficiency cost
 - Laziness or mistake of implementor
 - ...

Guaranteeing Uniqueness of Nonce



- Issues with nonce generation:
 - Counter needs storage
 - Need synchronization or transmission
 - Efficiency cost
 - Laziness or mistake of implementor
 - ...
- Sometimes, attacker can use same nonce multiple times

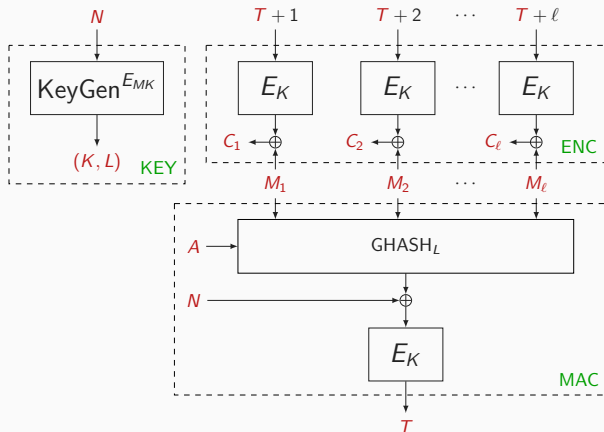
Nonce-Disrespecting Adversaries: Practical Forgery Attacks on GCM in TLS

Böck et al., USENIX WOOT 2016

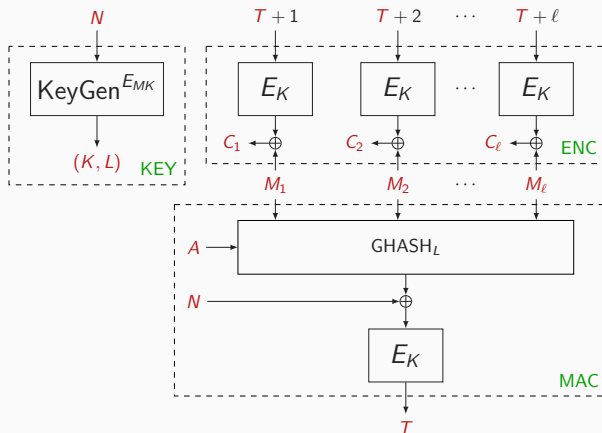
- GCM is widely used authenticated encryption scheme
- Used in TLS (“https”)
- Internet-wide scan for GCM implementations
- 184 devices with duplicated nonces
 - VISA, Polish bank, German stock exchange, ...
- ≈ 70.000 devices with random nonce

Intuition

- All input should be cryptographically transformed
- Any change in $(N, A, M) \longrightarrow$ unpredictable (C, T)
- Often comes at a price:
 - Efficiency
 - Security
 - Parallelizability
 - ...




- Gueron and Lindell (2015)
- MtE design
- Ongoing standardization (IETF RFC)
- Patent-free



- Gueron and Lindell (2015)
- MtE design
- Ongoing standardization (IETF RFC)
- Patent-free
- Inherits GCM features
- Secure against nonce-reuse
- Proof: Iwata and Seurin (2017)


Tweakable Block Ciphers

Tweakable Block Ciphers




- Using key K , message M is bijectively transformed to ciphertext C

Tweakable Block Ciphers




- Using key K , message M is bijectively transformed to ciphertext C
- Tweak T : flexibility to the cipher
- Each tweak gives different permutation

Tweakable Block Ciphers




- Using key K , message M is bijectively transformed to ciphertext C
- Tweak T : flexibility to the cipher
- Each tweak gives different permutation
- A good tweakable block cipher should behave like a random tweakable permutation

Tweakable Block Cipher Security




- Two oracles: \tilde{E}_K (for secret key K) and \tilde{p} (secret)

Tweakable Block Cipher Security




- Two oracles: \tilde{E}_K (for secret key K) and \tilde{p} (secret)
- Distinguisher \mathcal{D} has query access to one of these

Tweakable Block Cipher Security



- Two oracles: \tilde{E}_K (for secret key K) and \tilde{p} (secret)
- Distinguisher \mathcal{D} has query access to one of these
- \mathcal{D} tries to determine which oracle it communicates with


Tweakable Block Cipher Security



- Two oracles: \tilde{E}_K (for secret key K) and \tilde{p} (secret)
- Distinguisher \mathcal{D} has query access to one of these
- \mathcal{D} tries to determine which oracle it communicates with
- Its advantage is defined as:

$$\mathbf{Adv}_{\tilde{E}}^{\text{tprp}}(\mathcal{D}) = \Delta_{\mathcal{D}}(\tilde{E}_K ; \tilde{p}) = \left| \Pr(\mathcal{D}^{\tilde{E}_K} = 1) - \Pr(\mathcal{D}^{\tilde{p}} = 1) \right|$$

Tweakable Block Cipher Security




- Two oracles: \tilde{E}_K (for secret key K) and \tilde{p} (secret)
- Distinguisher \mathcal{D} has query access to one of these
- \mathcal{D} tries to determine which oracle it communicates with
- Its advantage is defined as:

$$\mathbf{Adv}_{\tilde{E}}^{\text{tprp}}(\mathcal{D}) = \Delta_{\mathcal{D}}(\tilde{E}_K ; \tilde{p}) = \left| \Pr(\mathcal{D}^{\tilde{E}_K} = 1) - \Pr(\mathcal{D}^{\tilde{p}} = 1) \right|$$


- $\mathbf{Adv}_{\tilde{E}}^{\text{tprp}}(q)$: supremal advantage over any \mathcal{D} with query complexity q

Strong Tweakable Block Cipher Security




- Two oracles: $(\tilde{E}_K, \tilde{E}_K^{-1})$ (for secret key K) and $(\tilde{p}, \tilde{p}^{-1})$ (secret)

Strong Tweakable Block Cipher Security




- Two oracles: $(\tilde{E}_K, \tilde{E}_K^{-1})$ (for secret key K) and $(\tilde{p}, \tilde{p}^{-1})$ (secret)
- Distinguisher \mathcal{D} has query access to one of these

Strong Tweakable Block Cipher Security



- Two oracles: $(\tilde{E}_K, \tilde{E}_K^{-1})$ (for secret key K) and $(\tilde{p}, \tilde{p}^{-1})$ (secret)
- Distinguisher \mathcal{D} has query access to one of these
- \mathcal{D} tries to determine which oracle it communicates with


Strong Tweakable Block Cipher Security



- Two oracles: $(\tilde{E}_K, \tilde{E}_K^{-1})$ (for secret key K) and $(\tilde{p}, \tilde{p}^{-1})$ (secret)
- Distinguisher \mathcal{D} has query access to one of these
- \mathcal{D} tries to determine which oracle it communicates with
- Its advantage is defined as:

$$\mathbf{Adv}_{\tilde{E}}^{\text{stprp}}(\mathcal{D}) = \Delta_{\mathcal{D}} \left(\tilde{E}_K, \tilde{E}_K^{-1} ; \tilde{p}, \tilde{p}^{-1} \right) = \left| \Pr \left(\mathcal{D}^{\tilde{E}_K, \tilde{E}_K^{-1}} = 1 \right) - \Pr \left(\mathcal{D}^{\tilde{p}, \tilde{p}^{-1}} = 1 \right) \right|$$

Strong Tweakable Block Cipher Security

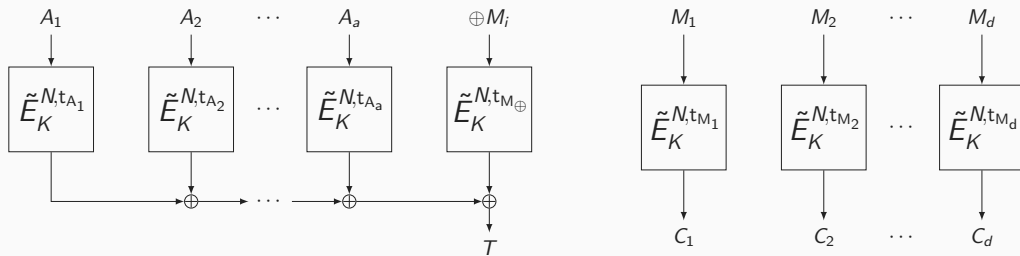


- Two oracles: $(\tilde{E}_K, \tilde{E}_K^{-1})$ (for secret key K) and $(\tilde{p}, \tilde{p}^{-1})$ (secret)
- Distinguisher \mathcal{D} has query access to one of these
- \mathcal{D} tries to determine which oracle it communicates with
- Its advantage is defined as:

$$\mathbf{Adv}_{\tilde{E}}^{\text{stprp}}(\mathcal{D}) = \Delta_{\mathcal{D}}(\tilde{E}_K, \tilde{E}_K^{-1}; \tilde{p}, \tilde{p}^{-1}) = \left| \Pr(\mathcal{D}^{\tilde{E}_K, \tilde{E}_K^{-1}} = 1) - \Pr(\mathcal{D}^{\tilde{p}, \tilde{p}^{-1}} = 1) \right|$$


- $\mathbf{Adv}_{\tilde{E}}^{\text{stprp}}(q)$: supremal advantage over any \mathcal{D} with query complexity q

Example Use in Θ CB (1/2)




- Generalized OCB by Rogaway et al. [RBBK01,Rog04,KR11]

Example Use in Θ CB (1/2)



- Generalized OCB by Rogaway et al. [RBBK01,Rog04,KR11]
- Internally based on tweakable block cipher \tilde{E}
 - Tweak (N, tweak) is unique for **every** evaluation
 - Different blocks always transformed under different tweak


Example Use in Θ CB (1/2)



- Generalized OCB by Rogaway et al. [RBBK01,Rog04,KR11]
- Internally based on tweakable block cipher \tilde{E}
 - Tweak (N, tweak) is unique for **every** evaluation
 - Different blocks always transformed under different tweak

$$\text{Adv}_{\text{AE}[\tilde{E}_k]}^{\text{ae}}(q)$$


Example Use in Θ CB (1/2)





- Generalized OCB by Rogaway et al. [RBBK01,Rog04,KR11]
- Internally based on tweakable block cipher \tilde{E}
 - Tweak (N, tweak) is unique for **every** evaluation
 - Different blocks always transformed under different tweak
- Triangle inequality:

$$\text{Adv}_{\text{AE}[\tilde{E}_k]}^{\text{ae}}(q) \leq \text{Adv}_{\text{AE}[\tilde{p}]}^{\text{ae}}(q) + \text{Adv}_{\tilde{E}}^{\text{stprp}}(q)$$

Example Use in Θ CB (2/2)




Example Use in Θ CB (2/2)




- Nonce uniqueness \Rightarrow tweak uniqueness

Example Use in Θ CB (2/2)




- Nonce uniqueness \Rightarrow tweak uniqueness
- Encryption calls behave like random functions: $AE[\tilde{p}] = \$$

Example Use in Θ CB (2/2)




- Nonce uniqueness \Rightarrow tweak uniqueness
- Encryption calls behave like random functions: $AE[\tilde{p}] = \$$
- Authentication almost behaves like random function (but nonces may repeat)

Example Use in Θ CB (2/2)



- Nonce uniqueness \Rightarrow tweak uniqueness
- Encryption calls behave like random functions: $AE[\tilde{p}] = \$$
- Authentication almost behaves like random function (but nonces may repeat)
 - Tag forged with probability at most $1/(2^n - 1)$

Example Use in Θ CB (2/2)




- Nonce uniqueness \Rightarrow tweak uniqueness
- Encryption calls behave like random functions: $AE[\tilde{p}] = \$$
- Authentication almost behaves like random function (but nonces may repeat)
 - Tag forged with probability at most $1/(2^n - 1)$

$$\mathbf{Adv}_{AE[\tilde{p}]}^{\text{ae}}(q) \leq 1/(2^n - 1)$$


Building Tweakable Block Ciphers

- TWEAKEY by Jean et al. [JNP14]:




- f : round function
- g : subkey computation
- h : transformation of (K, T)

- TWEAKEY by Jean et al. [JNP14]:




- f : round function
- g : subkey computation
- h : transformation of (K, T)
- Security measured through cryptanalysis
- Our focus: modular design

- LRW_1 and LRW_2 by Liskov et al. [LRW02]:




- h is XOR-universal hash
 - E.g., $h(T) = h \otimes T$ for n -bit “key” h

- XEX by Rogaway [Rog04]:



- $(\alpha, \beta, \gamma, N)$ is tweak (simplified)



- XEX by Rogaway [Rog04]:



- $(\alpha, \beta, \gamma, N)$ is tweak (simplified)
- Used in OCB2 and ± 14 CAESAR candidates

Powering-Up Masking (XEX)



- XEX by Rogaway [Rog04]:



- $(\alpha, \beta, \gamma, N)$ is tweak (simplified)
- Used in OCB2 and ± 14 CAESAR candidates
- Permutation-based variants in Minalpher and Prøst (generalized by Cogliati et al. [CLS15])

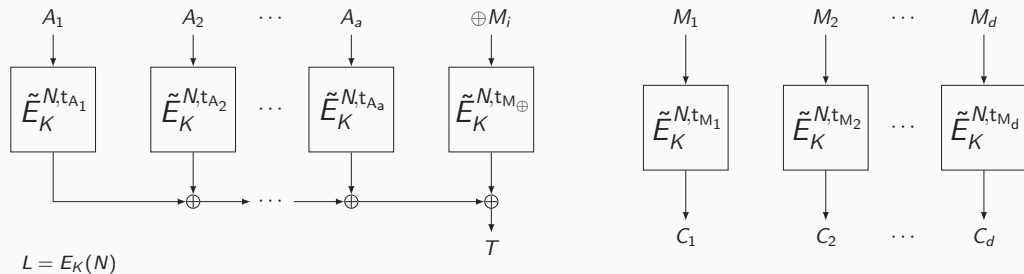
Powering-Up Masking (XEX)

- XEX by Rogaway [Rog04]:

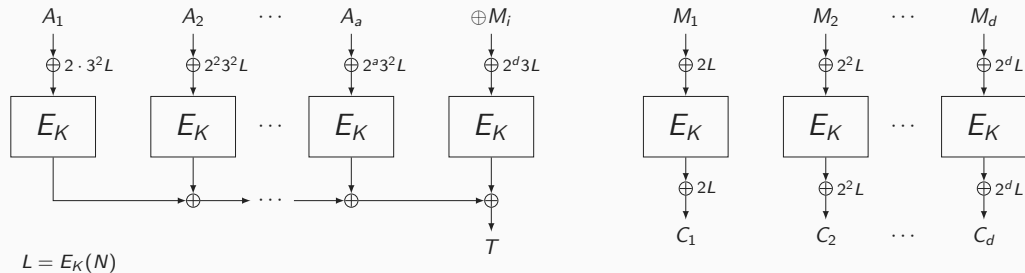


- $(\alpha, \beta, \gamma, N)$ is tweak (simplified)
- Used in OCB2 and ± 14 CAESAR candidates
- Permutation-based variants in Minalpher and Prøst (generalized by Cogliati et al. [CLS15])
- STPRP up to $2^{n/2}$ queries **provided masks are all distinct**

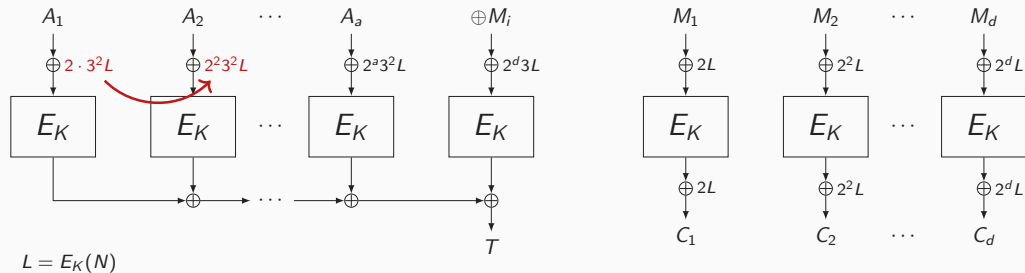
Powering-Up Masking in OCB2-Like Construction



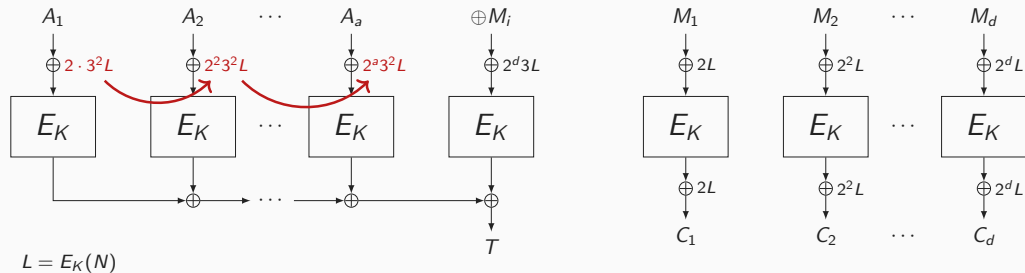
Powering-Up Masking in OCB2-Like Construction





Powering-Up Masking in OCB2-Like Construction



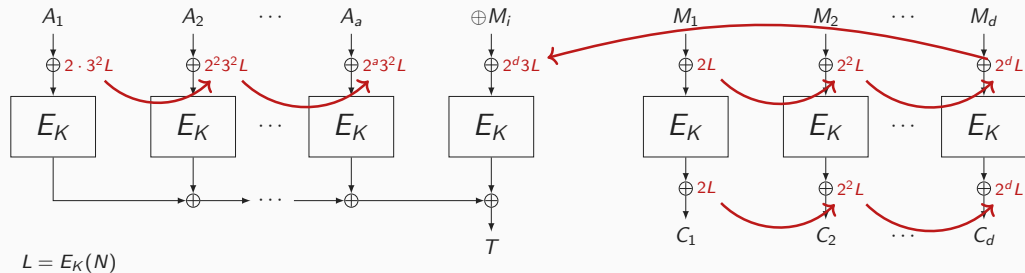
Powering-Up Masking in OCB2-Like Construction



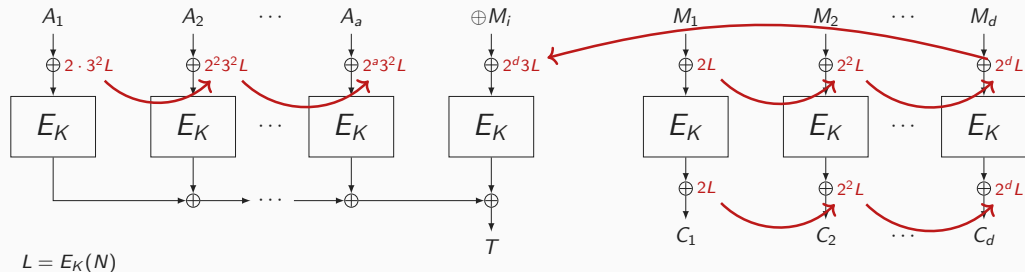
Powering-Up Masking in OCB2-Like Construction



Powering-Up Masking in OCB2-Like Construction

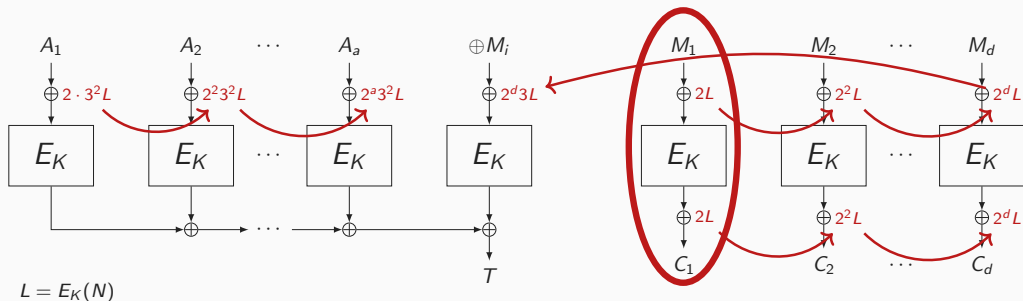


Powering-Up Masking in OCB2-Like Construction



- Update of mask:
 - Shift and conditional XOR
- Variable time computation
- Expensive on certain platforms


Intermezzo: Why Start at $2 \cdot E_K(N)$? (1/2)



- Update of mask:
 - Shift and conditional XOR
- Variable time computation
- Expensive on certain platforms


Intermezzo: Why Start at $2 \cdot E_K(N)$? (2/2)

- Suppose we would mask with $E_K(N)$:



Intermezzo: Why Start at $2 \cdot E_K(N)$? (2/2)


- Suppose we would mask with $E_K(N)$:



- Distinguisher can make inverse queries

Intermezzo: Why Start at $2 \cdot E_K(N)$? (2/2)


- Suppose we would mask with $E_K(N)$:



- Distinguisher can make inverse queries
- Putting $C = 0$ gives $M = N \oplus E_K(N)$


Intermezzo: Why Start at $2 \cdot E_K(N)$? (2/2)

- Suppose we would mask with $E_K(N)$:




- Distinguisher can make inverse queries
- Putting $C = 0$ gives $M = N \oplus E_K(N)$
- Distinguisher knows N so learns “subkey” $E_K(N)$

- XEX by Rogaway [Rog04]:




- $(\alpha, \beta, \gamma, N)$ is tweak (simplified)

- XEX by Rogaway [Rog04]:




- $(\alpha, \beta, \gamma, N)$ is tweak (simplified)
- (α, β, γ) must be from a certain **admissible domain**
- We need that $2^\alpha 3^\beta 7^\gamma \neq 2^{\alpha'} 3^{\beta'} 7^{\gamma'}$ for any $(\alpha, \beta, \gamma) \neq (\alpha', \beta', \gamma')$
 - Otherwise, attacker can obviously break the scheme

- XEX by Rogaway [Rog04]:




- $(\alpha, \beta, \gamma, N)$ is tweak (simplified)
- (α, β, γ) must be from a certain **admissible domain**
- We need that $2^\alpha 3^\beta 7^\gamma \neq 2^{\alpha'} 3^{\beta'} 7^{\gamma'}$ for any $(\alpha, \beta, \gamma) \neq (\alpha', \beta', \gamma')$
 - Otherwise, attacker can obviously break the scheme
- Typical: $\alpha \in \{1, \dots, \text{large}\}$, and $\beta, \gamma \in \{0, 1, 2\}$

Intermezzo: XTS Disk Encryption Mode (1/2)




- XTS = XEX-based Tweaked-codebook mode with ciphertext Stealing

Intermezzo: XTS Disk Encryption Mode (1/2)




- XTS = XEX-based Tweaked-codebook mode with ciphertext Stealing
 - Electronic CodeBook (ECB) ...

Intermezzo: XTS Disk Encryption Mode (1/2)




- XTS = XEX-based Tweaked-codebook mode with ciphertext Stealing
 - Electronic CodeBook (ECB) ...
 - ... with XEX as primitive (i is sector, j is block within sector) ...

Intermezzo: XTS Disk Encryption Mode (1/2)




- XTS = XEX-based Tweaked-codebook mode with ciphertext Stealing
 - Electronic CodeBook (ECB) ...
 - ...with XEX as primitive (i is sector, j is block within sector) ...
 - ...and doing a fancy thing called “ciphertext stealing”

Intermezzo: XTS Disk Encryption Mode (1/2)




- XTS = XEX-based Tweaked-codebook mode with ciphertext Stealing
 - Electronic CodeBook (ECB) ...
 - ...with XEX as primitive (i is sector, j is block within sector) ...
 - ...and doing a fancy thing called “ciphertext stealing”

Intermezzo: XTS Disk Encryption Mode (1/2)




- XTS = XEX-based Tweaked-codebook mode with ciphertext Stealing
 - Electronic CodeBook (ECB) ...
 - ...with XEX as primitive (i is sector, j is block within sector) ...
 - ...and doing a fancy thing called “ciphertext stealing”
- One sector consists of 512 bytes, or 32 blocks

Intermezzo: XTS Disk Encryption Mode (2/2)




- Features:
 - Tweak **unique** for every block, changing tweak is **efficient**
 - Incrementality: change in one (or few) blocks

Intermezzo: XTS Disk Encryption Mode (2/2)




- Features:
 - Tweak **unique** for every block, changing tweak is **efficient**
 - Incrementality: change in one (or few) blocks
- XTS-AES is standardized as IEEE P1619
- Supported by myriad disk encryption tools: BestCrypt, dm-crypt, TrueCrypt, VeraCrypt, DiskCryptor, FileVault 2 (MacOS), BitLocker (Windows 10)

- OCB1 and OCB3 use Gray Codes:




- (α, N) is tweak
- Updating: $G(\alpha) = G(\alpha - 1) \oplus 2^{\text{ntz}(\alpha)}$

- OCB1 and OCB3 use Gray Codes:




- (α, N) is tweak
- Updating: $G(\alpha) = G(\alpha - 1) \oplus 2^{\text{ntz}(\alpha)}$
 - Single XOR
 - Logarithmic amount of field doublings (precomputed)
- More efficient than powering-up [KR11]

- MEM by Granger et al. [GJMN16]:




- φ_i are fixed LFSRs, $(\alpha, \beta, \gamma, N)$ is tweak (simplified)

- MEM by Granger et al. [GJMN16]:




- φ_i are fixed LFSRs, $(\alpha, \beta, \gamma, N)$ is tweak (simplified)
- Combines advantages of:
 - Powering-up masking
 - Word-based LFSRs

- MEM by Granger et al. [GJMN16]:



- φ_i are fixed LFSRs, $(\alpha, \beta, \gamma, N)$ is tweak (simplified)
- Combines advantages of:
 - Powering-up masking
 - Word-based LFSRs
- Simpler, constant-time (by default), more efficient

Application to AE: OPP



$$L = P(N \| k)$$

$$\varphi_1 = \varphi \oplus id, \varphi_2 = \varphi^2 \oplus \varphi \oplus id$$

- Offset Public Permutation (OPP)
- Generalization of OCB3:
 - Permutation-based
 - More efficient MEM masking
- Security against nonce-respecting adversaries

- US NIST recently currently ran competition for lightweight cryptography
- Round 1: 56 submissions in February 2019
- Round 2: 32 submissions in August 2019
- Final round: 10 submissions in March 2021
- Winner **Ascon** announced February 2023

- US NIST recently currently ran competition for **lightweight cryptography**
 - Round 1: 56 submissions in February 2019
 - Round 2: 32 submissions in August 2019
 - Final round: 10 submissions in March 2021
 - Winner **Ascon** announced February 2023
-
- Some submissions were sponge-based (like Ascon)
 - Some submissions used techniques from this lecture