

Cryptographic Hash Functions and Key Derivation

Applied Cryptography – Spring 2024

Bart Mennink

February 26, 2024

Institute for Computing and Information Sciences
Radboud University

1 / 35

Last Lectures

- We learned the basics of symmetric cryptography:
 - Encryption
 - Message authentication
 - Authenticated encryption
- These can be built from (a.o.):
 - Tweakable block ciphers
 - Block ciphers
 - Permutations
- There is one more core functionality of symmetric cryptography:

Cryptographic hashing

2 / 35

Outline

Hash Functions

History

Indifferentiability

Sponges

Keccak and SHA-3


Key Derivation Functions

Conclusion

3 / 35

Hash Functions

Hash Functions



- Function \mathcal{H} from $\{0, 1\}^*$ to $\{0, 1\}^n$
 - No key input
 - Variable-length input
 - Classically fixed length output (but could be variable as well)

4 / 35

Example: Digital Signatures

- Suppose you want to sign a message M with a private key PrK :

$$\sigma = \text{sign}(PrK, M)$$

- You can send (M, σ) to the receiver
- The receiver can use your public key PK to verify:

$$\text{verify}(PK, M, \sigma)$$

- If M is huge, computing $\text{sign}(PrK, M)$ can be costly
- One solution is to sign $\mathcal{H}(M)$ instead: $\sigma = \text{sign}(PrK, \mathcal{H}(M))$
- This is fine, as long as one cannot come up with two different messages M, M' that hash to the same value!
- This is called **collision resistance** of the hash function

5 / 35

Example: Forging Digital Signatures

- Sometimes, **collision resistance** is a too strong requirement
- Suppose you intercept a message M with a signature $\sigma = \text{sign}(PrK, \mathcal{H}(M))$
- A forgery would be a different message M' with

$$\sigma = \text{sign}(PrK, \mathcal{H}(M'))$$

- For this, it is *sufficient* to find a message M' that has the same hash value as M
- So we require it to be hard for an attacker to find such a message for \mathcal{H}
- This is called **second preimage resistance**

6 / 35

Example: Password Hashing

- Consider a server that stores hashes of passwords: $\text{hash} = \mathcal{H}(\text{password}, \text{salt})$
 - Authentication is done by entering **password** and verifying **hash**
- Suppose an adversary gets possession of **hash** and **salt**
- It manages to **pass authentication** if it can find a **password** for the given **hash** and **salt**
- So we require it to be hard for an attacker to find such a preimage for \mathcal{H}
- This is called **preimage resistance**

7 / 35

Further Examples and Security Requirements

Many More Applications of Hash Functions

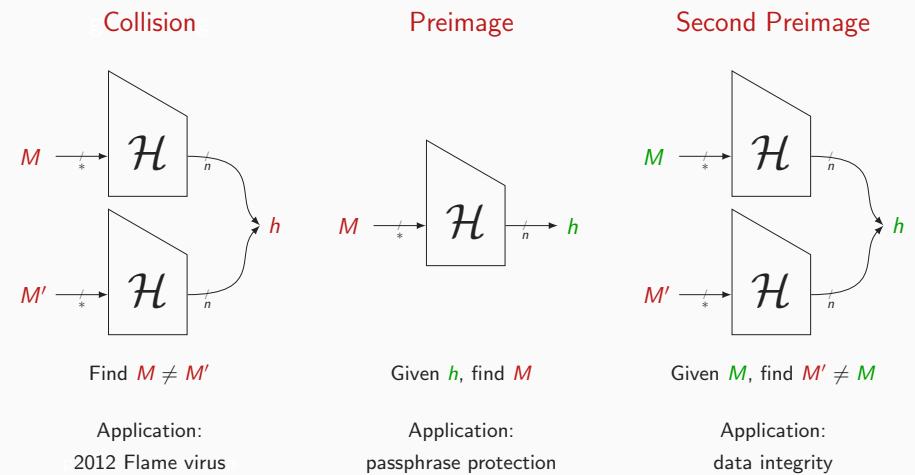
- Destroying algebraic structure, e.g.,
 - Encryption with RSA: OAEP
 - Signing with RSA: PSS

Security Model?

- Expressing security model is not easy
- We have seen examples of **collision**, **preimage**, and **second preimage resistance**
 - These are the **classical** security requirements
 - Focal point of first part of lecture
- Ideally, we want that a hash function behaves like a **RO**
 - This is theoretically **impossible**
 - A security model that still solves this somewhat, is **indifferentiability**
 - Focal point of second part of lecture

8 / 35


Classical Security Requirements



9 / 35

History

Hash Functions from Compression Functions (1/2)




Merkle-Damgård with Strengthening

- Damgård [Dam89] and Merkle [Mer89]
- Consecutive evaluation of compression function F
- Length encoding at the end
- Used in MD5, SHA-1/2, ...
- Not a very good scheme, as we will see

10 / 35

Hash Functions from Compression Functions (2/2)



Security of Merkle-Damgård

- \mathcal{H} and F have same security models
- We happen to have (up to some degree):

$$F \text{ is col/sec/pre secure} \implies \mathcal{H} \text{ is col/sec/pre secure}$$

11 / 35


MD5 and NIST Standards SHA-1/2

- MD5 [Rivest, 1991]
 - Based on MD4 that was an original design
 - 128-bit digest
- SHA-1 [NIST, 1995] (after SHA-0 [NIST, 1993])
 - Inspired by MD5, designed at NSA
 - 160-bit digest
- SHA-2 series [NIST, 2001/2008]
 - *Reinforced versions of SHA-1*, designed at NSA
 - 6 functions with 224-, 256-, 384- and 512-bit digest
- Internally (for each of these):
 - Merkle-Damgård iteration mode
 - F based on a block cipher E in Davies-Meyer mode
 - Block cipher E : software oriented word-based design

12 / 35

Indifferentiability


Indistinguishability of Hash Functions (1/3)



- \mathcal{H} should **behave like** random oracle \mathcal{RO}
- But \mathcal{H} is **not a random system**
 - Distinguisher can distinguish \mathcal{H} from \mathcal{RO} with probability 1
- Solution: **introduce randomness**

13 / 35


Indistinguishability of Hash Functions (2/3)



- \mathcal{H}^P for random primitive \mathcal{P} should behave like random oracle \mathcal{RO}
- \mathcal{P} can be ideal function F , block cipher E , permutation P , ...
- Adversarial model still too weak, we don't want to base security on secrecy of \mathcal{P}
- Distinguisher should be able to evaluate \mathcal{P}
- Solution: give \mathcal{D} access to \mathcal{P}

14 / 35


Indistinguishability of Hash Functions (3/3)



- $(\mathcal{H}^P, \mathcal{P})$ for random primitive \mathcal{P} should behave like random oracle $(\mathcal{RO}, \mathcal{P})$
- Adversary can still trivially distinguish:
 - Make a single construction query (to \mathcal{H}^P or \mathcal{RO})
 - Simulate \mathcal{H}^P using the oracle \mathcal{P}
- In the real world, the responses are consistent, in the ideal world they are not
- Solution: **indifferentiability**

15 / 35


Indifferentiability (1/2)



- Maurer et al. [MRH04] and Coron et al. [CDMP05]
- $(\mathcal{H}^P, \mathcal{P})$ for random primitive \mathcal{P} should behave like random oracle \mathcal{RO} paired with a simulator \mathcal{S} that maintains construction-primitive consistency
- Based on **composition**: distinguisher in one game is simulator in another one

16 / 35

Indifferentiability (2/2)



- \mathcal{H} is **indifferentiable** from \mathcal{RO} if for some simulator \mathcal{S} :
 $\Delta_{\mathcal{D}}(\mathcal{H}, \mathcal{P}; \mathcal{RO}, \mathcal{S})$ is small
- Proof idea:
 - Step 1.** Construct a **clever** simulator \mathcal{S}
 - Step 2.** Use game-playing or H-coefficient technique (not included in course)
- Unfortunately, proofs are often very tedious
- Indifferentiability \implies coll/pre/sec security

17 / 35

Differentiability of Merkle-Damgård (1/2)

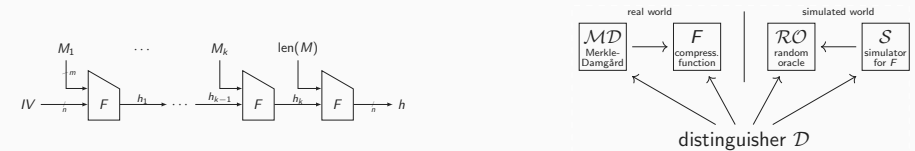


Merkle-Damgård is Easily Differentiable

- Goal is to prove that there **exists** a distinguisher that fools **any** simulator
- Let \mathcal{S} be **any** simulator
- Denote construction oracle by $\mathcal{H} \in \{\mathcal{MD}, \mathcal{RO}\}$ and primitive by $\mathcal{P} \in \{F, \mathcal{S}\}$

18 / 35

Differentiability of Merkle-Damgård (2/2)



Merkle-Damgård is Easily Differentiable

- Distinguisher \mathcal{D} operates as follows:
 - Pick arbitrary M_1
 - Query $\mathcal{H}(M_1) = h$, $\mathcal{H}(M_1 \parallel \text{len}(M_1)) = h'$, and $\mathcal{P}(h, \text{len}(M_1 \parallel \text{len}(M_1))) = y$
 - Verify if $h' \stackrel{?}{=} y$
- Real world: $h' = y$ by design
- Simulated world: \mathcal{S} must choose output y only based on knowledge of h and $\text{len}(M_1 \parallel \text{len}(M_1))$, but it cannot deduce M_1 from these values and it will likely fail

19 / 35

Sponges

A Bit of History (1/3)

- 2005-2006: MD5 and SHA-1 crisis
 - Actual collisions for MD5
 - Theoretical collision attacks for SHA-1
 - Attacks on Merkle-Damgård with higher success probability than believed up to that point
- SHA-2 based on same principles, so US NIST got nervous
- 2007: NIST announces plans to have open SHA-3 competition
 - Goal: find a worthy successor for SHA-2
 - Similar process as AES competition
- 2008: NIST publishes SHA-3 requirements
 - *More efficient than SHA-2*
 - Output lengths: 224, 256, 384, 512 bits
 - Security: collision and (second) preimage resistance

Slide credit: Joan Daemen

20 / 35

A Bit of History (2/3)

- Competition started in 2008
- Three-round public process
 - round 1: 64 submissions, 51 accepted
 - round 2: 14 semi-finalists
 - round 3: 5 finalists
- All selections done by NIST but based on public evaluation by crypto community
- October 2012: NIST announces the SHA-3 winner
- The winner: **Keccak**
 - By Guido Bertoni, Joan Daemen, Michaël Peeters and Gilles Van Assche
 - Something completely different than MD5/SHA-1/SHA-2 ...
 - ... and completely different than Rijndael/AES
- August 2015: NIST finally publishes the SHA-3 standard: FIPS 202

Slide credit: Joan Daemen

21 / 35

A Bit of History (3/3)

- Keccak is a permutation-based hash function, a **sponge**
- Sponge differs from Merkle-Damgård in two main ways

1. Merkle-Damgård Functions Designed With Property Preservation in Mind


- F must be collision resistant for \mathcal{H} to be collision resistant
- But this means: we require F to be cryptographically strong
- This often incurs efficiency penalty
- **Solution in sponge:** skip reduction step and get cleaner and more efficient design

2. Block Ciphers Have a Key Schedule and Data Path

- F is in turn often built from a block cipher (like Davies-Meyer)
- While data paths are reasonably well-understood, key schedules not so much
- In addition, final state of key schedule is discarded
- Block cipher is weirdly compressing function from $n + k$ to n bits
- **Solution in sponge:** use (iterative) permutation from b to b bits

22 / 35


Ancient Definition of Hashing



- Function \mathcal{H} from $\{0, 1\}^*$ to $\{0, 1\}^n$
 - Variable-length input
 - Fixed-length output
 - Mode *on top of* \mathcal{H} might give variable-length output

23 / 35


Modern Definition of Hashing



- Function \mathcal{XOF} from $\{0, 1\}^*$ to $\{0, 1\}^\infty$
 - Variable-length input
 - Variable-length output
 - User specifies output length n when calling the function

24 / 35

Sponges [BDPV07]



- P is a b -bit permutation, with $b = r + c$
 - r is the rate
 - c is the capacity (security parameter)

25 / 35

Indifferentiability of the Sponge [BDPV08]

- Assume that P is a random permutation
- Sponge indifferentiable from RO: $\Delta_{\mathcal{D}}(\text{Sponge}, P; \mathcal{RO}, \mathcal{S}) \leq N^2/2^{c+1}$
 - N is number of permutation evaluations that attacker can make
 - Collisions in the inner part break security of the sponge
- Security of sponge truncated to n bits against classical attacks:

Collision resistance: $N^2/2^{c+1} + N^2/2^{n+1}$

Preimage resistance: $N^2/2^{c+1} + N/2^n$


Second preimage resistance: $N^2/2^{c+1} + N/2^n$

distance from sponge to RO
(N is # primitive evaluations)

classical attacks against RO
(N is # oracle evaluations)

26 / 35

Sponge Recap




- Relevant parameters:
 - c : capacity – typically twice the security strength
 - r : rate – amount of bits absorbed/squeezed per permutation
 - b : width of permutation – $b = r + c$
 - n : amount of output bits
- Security strength (for random sponge):
 - collision resistance: $\min(c/2, n/2)$
 - first and second preimage resistance: $\min(c/2, n)$

27 / 35

Keccak and SHA-3

Keccak and Keccak-f


- Keccak is a sponge function using permutation Keccak-f
- Keccak-f operates on 3-dimensional state:
 - 5×5 lanes, each containing 2^ℓ bits (1, 2, 4, 8, 16, 32 or 64)
 - (5×5) -bit slices, 2^ℓ of them



Slide credit: Joan Daemen

28 / 35

Keccak-f: Steps of the Round Function



bit-oriented highly-symmetric *wide-trail* design

Slide credit: Joan Daemen

29 / 35

Keccak[r, c]


- Keccak[r, c] is a sponge function using permutation Keccak-f
 - 7 permutations: $b \in \{25, 50, 100, 200, 400, 800, 1600\}$ from toy over lightweight to high-speed
- SHA-3 instance SHAKE128: $r = 1344$ and $c = 256$
 - Permutation width: 1600
 - Security strength: 128
- Lightweight instance: $r = 40$ and $c = 160$
 - Permutation width: 200
 - Security strength: 80 (what SHA-1 should have offered)
- Security status:
 - Best attack on hash function covers 6-round version
 - # rounds ranges from 18 for $b = 200$ to 24 for $b = 1600$

Slide credit: Joan Daemen

30 / 35

Key Derivation Functions

Key Derivation Functions



- Derive secret key from a password, passphrase, ...
- Key stretching, strengthening, ...
- Key diversification
- ...

31 / 35

Intermezzo: HMAC Message Authentication Code

How to Build Hash-Based PRF?


- Ideally, one does $\text{PRF}(K, M) = \mathcal{H}(K \parallel M)$
- For the sponge, that works (why?) (more about this next week)
- For ancient hash functions, like SHA-1 and SHA-2, this does not work (why?)
- Still, many people use these functions, and, sponges are “quite” recent
- People searched for “inventive” ways to turn a hash function into a PRF

HMAC (Bellare et al. [BCK96])

- Let opad be a constant string consisting of repetition of 0x5c
 ipad be a constant string consisting of repetition of 0x36
- $\text{HMAC}(K \parallel M) = \mathcal{H}(K \oplus \text{opad} \parallel \mathcal{H}(K \oplus \text{ipad} \parallel M))$
- **Band-aid cryptography**, not the most beautiful construction, but **very popular!**

32 / 35

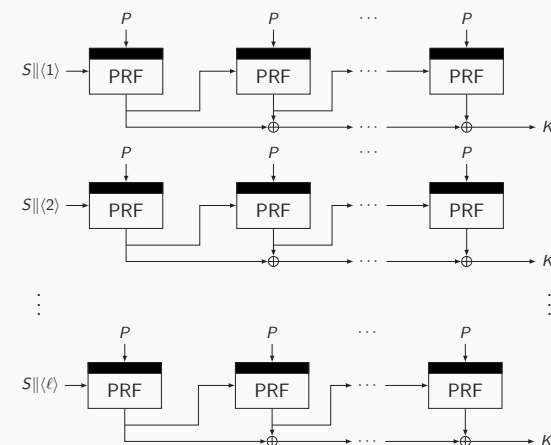
HKDF Key Derivation Function



- RFC 5869 (2010)
- “info” is optional material, e.g., to bind application to use case

33 / 35

PBKDF2 Key Derivation Function



- RFC 2898 (2000)
- Standardized in PKCS #5 v2.0
- Popular PRF choices:
 - HMAC-SHA-1 (in WPA2)
 - HMAC-SHA-256, HMAC-SHA-512, HMAC-RIPEMD-160 (in VeraCrypt)

34 / 35

Conclusion

Next Week

- Sponge construction solved the problems that were present in Merkle-Damgård
- No band-aid-type cryptography (like HMAC) needed
 - $\text{PRF}(K, M) = \text{sponge}(K \parallel M)$ would have done the job
- Sponges can also be used for
 - Message authentication
 - Keystream generation
 - Authenticated encryption
 - ...
- This will be the topic of next week!